

ass-4-lokesh

September 21, 2023

```
[1]: import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
```

```
[2]: df = pd.read_csv('/content/winequality-red.csv')
df.head()
```

```
[2]:
```

	fixed acidity	volatile acidity	citric acid	residual sugar	chlorides	\
0	7.4	0.70	0.00	1.9	0.076	
1	7.8	0.88	0.00	2.6	0.098	
2	7.8	0.76	0.04	2.3	0.092	
3	11.2	0.28	0.56	1.9	0.075	
4	7.4	0.70	0.00	1.9	0.076	

	free sulfur dioxide	total sulfur dioxide	density	pH	sulphates	\
0	11.0	34.0	0.9978	3.51	0.56	
1	25.0	67.0	0.9968	3.20	0.68	
2	15.0	54.0	0.9970	3.26	0.65	
3	17.0	60.0	0.9980	3.16	0.58	
4	11.0	34.0	0.9978	3.51	0.56	

	alcohol	quality
0	9.4	5
1	9.8	5
2	9.8	5
3	9.8	6
4	9.4	5

```
[3]: df.describe()
```

```
[3]:
```

	fixed acidity	volatile acidity	citric acid	residual sugar	\
count	1599.000000	1599.000000	1599.000000	1599.000000	
mean	8.319637	0.527821	0.270976	2.538806	
std	1.741096	0.179060	0.194801	1.409928	
min	4.600000	0.120000	0.000000	0.900000	
25%	7.100000	0.390000	0.090000	1.900000	


50%	7.900000	0.520000	0.260000	2.200000
75%	9.200000	0.640000	0.420000	2.600000
max	15.900000	1.580000	1.000000	15.500000

	chlorides	free sulfur dioxide	total sulfur dioxide	density \
count	1599.000000	1599.000000	1599.000000	1599.000000
mean	0.087467	15.874922	46.467792	0.996747
std	0.047065	10.460157	32.895324	0.001887
min	0.012000	1.000000	6.000000	0.990070
25%	0.070000	7.000000	22.000000	0.995600
50%	0.079000	14.000000	38.000000	0.996750
75%	0.090000	21.000000	62.000000	0.997835
max	0.611000	72.000000	289.000000	1.003690

	pH	sulphates	alcohol	quality
count	1599.000000	1599.000000	1599.000000	1599.000000
mean	3.311113	0.658149	10.422983	5.636023
std	0.154386	0.169507	1.065668	0.807569
min	2.740000	0.330000	8.400000	3.000000
25%	3.210000	0.550000	9.500000	5.000000
50%	3.310000	0.620000	10.200000	6.000000
75%	3.400000	0.730000	11.100000	6.000000
max	4.010000	2.000000	14.900000	8.000000


```
[4]: plt.bar(df.chlorides.value_counts(),df.chlorides.value_counts())
```

```
[4]: <BarContainer object of 153 artists>
```



```
[8]: sns.jointplot(x="density",y="quality",data=df)
```

```
[8]: <seaborn.axisgrid.JointGrid at 0x7e6e96f39c30>
```



```
[9]: df.corr()
```

```
[9]:
```

	fixed acidity	volatile acidity	citric acid	\
fixed acidity	1.000000	-0.256131	0.671703	
volatile acidity	-0.256131	1.000000	-0.552496	
citric acid	0.671703	-0.552496	1.000000	
residual sugar	0.114777	0.001918	0.143577	
chlorides	0.093705	0.061298	0.203823	
free sulfur dioxide	-0.153794	-0.010504	-0.060978	
total sulfur dioxide	-0.113181	0.076470	0.035533	
density	0.668047	0.022026	0.364947	
pH	-0.682978	0.234937	-0.541904	
sulphates	0.183006	-0.260987	0.312770	

alcohol	-0.061668	-0.202288	0.109903
quality	0.124052	-0.390558	0.226373

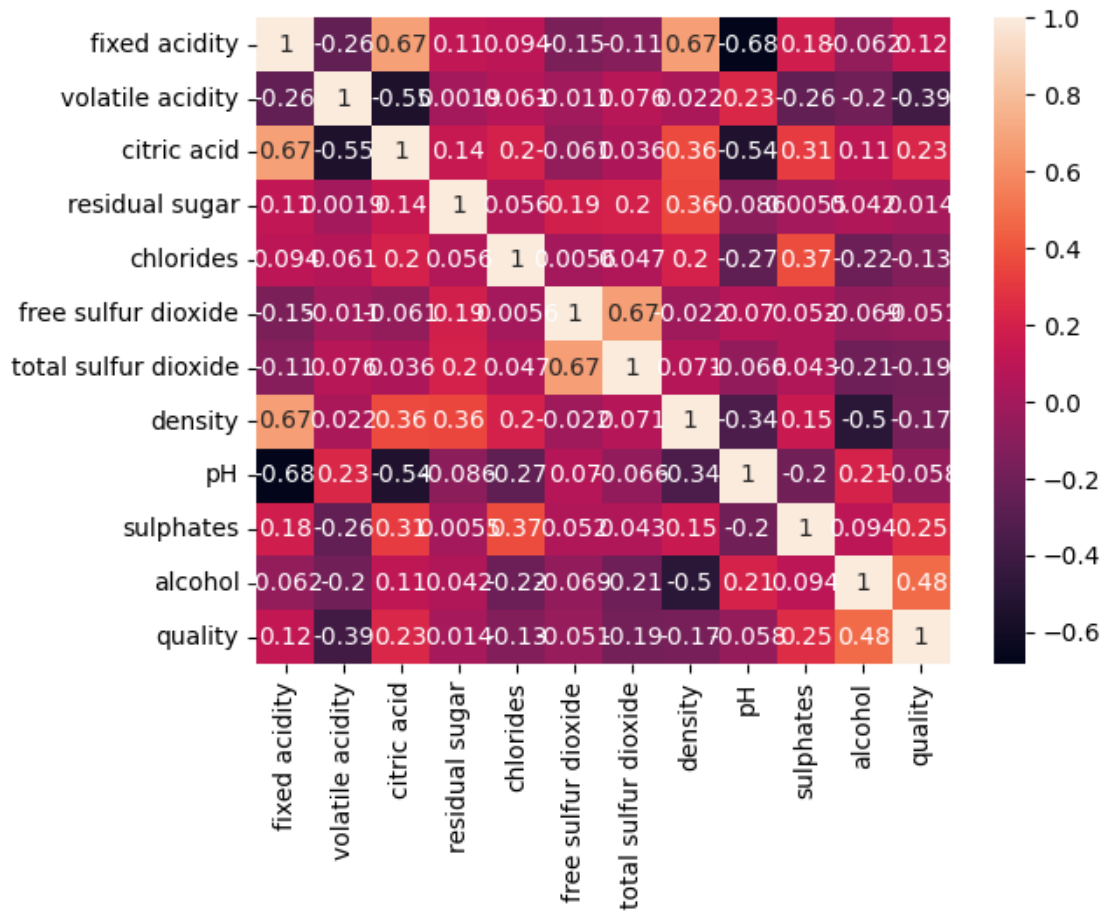
	residual sugar	chlorides	free sulfur dioxide	\
fixed acidity	0.114777	0.093705	-0.153794	
volatile acidity	0.001918	0.061298	-0.010504	
citric acid	0.143577	0.203823	-0.060978	
residual sugar	1.000000	0.055610	0.187049	
chlorides	0.055610	1.000000	0.005562	
free sulfur dioxide	0.187049	0.005562	1.000000	
total sulfur dioxide	0.203028	0.047400	0.667666	
density	0.355283	0.200632	-0.021946	
pH	-0.085652	-0.265026	0.070377	
sulphates	0.005527	0.371260	0.051658	
alcohol	0.042075	-0.221141	-0.069408	
quality	0.013732	-0.128907	-0.050656	

	total sulfur dioxide	density	pH	sulphates	\
fixed acidity	-0.113181	0.668047	-0.682978	0.183006	
volatile acidity	0.076470	0.022026	0.234937	-0.260987	
citric acid	0.035533	0.364947	-0.541904	0.312770	
residual sugar	0.203028	0.355283	-0.085652	0.005527	
chlorides	0.047400	0.200632	-0.265026	0.371260	
free sulfur dioxide	0.667666	-0.021946	0.070377	0.051658	
total sulfur dioxide	1.000000	0.071269	-0.066495	0.042947	
density	0.071269	1.000000	-0.341699	0.148506	
pH	-0.066495	-0.341699	1.000000	-0.196648	
sulphates	0.042947	0.148506	-0.196648	1.000000	
alcohol	-0.205654	-0.496180	0.205633	0.093595	
quality	-0.185100	-0.174919	-0.057731	0.251397	

	alcohol	quality
fixed acidity	-0.061668	0.124052
volatile acidity	-0.202288	-0.390558
citric acid	0.109903	0.226373
residual sugar	0.042075	0.013732
chlorides	-0.221141	-0.128907
free sulfur dioxide	-0.069408	-0.050656
total sulfur dioxide	-0.205654	-0.185100
density	-0.496180	-0.174919
pH	0.205633	-0.057731
sulphates	0.093595	0.251397
alcohol	1.000000	0.476166
quality	0.476166	1.000000

```
[10]: sns.heatmap(df.corr(),annot=True)
```

[10]: <Axes: >



```
[11]: from sklearn.model_selection import train_test_split
from sklearn.preprocessing import StandardScaler
X=df.drop("quality",axis=1)
y=df["quality"]
```

```
[13]: X.head()
```

```
[13]:    fixed acidity  volatile acidity  citric acid  residual sugar  chlorides \
0           7.4           0.70           0.00           1.9       0.076
1           7.8           0.88           0.00           2.6       0.098
2           7.8           0.76           0.04           2.3       0.092
3          11.2           0.28           0.56           1.9       0.075
4           7.4           0.70           0.00           1.9       0.076

    free sulfur dioxide  total sulfur dioxide  density  pH  sulphates \
0           11.0           34.0  0.9978  3.51       0.56
```

1	25.0	67.0	0.9968	3.20	0.68
2	15.0	54.0	0.9970	3.26	0.65
3	17.0	60.0	0.9980	3.16	0.58
4	11.0	34.0	0.9978	3.51	0.56

	alcohol
0	9.4
1	9.8
2	9.8
3	9.8
4	9.4

```
[16]: sc=StandardScaler()
X_scaled=sc.fit_transform(X)
X_train,X_test,y_train,y_test=train_test_split(X_scaled,y,test_size=0.
↪2,random_state=42)
```

```
[19]: from sklearn.linear_model import LogisticRegression
from sklearn.metrics import confusion_matrix
le = LogisticRegression()
```

```
[22]: model=le.fit(X_train,y_train)
y_pred = model.predict(X_test)
cm=confusion_matrix(y_test,y_pred)
```

```
/usr/local/lib/python3.10/dist-packages/sklearn/linear_model/_logistic.py:458:
ConvergenceWarning: lbfgs failed to converge (status=1):
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.
```

Increase the number of iterations (max_iter) or scale the data as shown in:

<https://scikit-learn.org/stable/modules/preprocessing.html>

Please also refer to the documentation for alternative solver options:

https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression

```
n_iter_i = _check_optimize_result(
```

```
[23]: from sklearn.metrics import accuracy_score,
↪confusion_matrix,classification_report,roc_auc_score,roc_curve
accuracy_score(y_test,y_pred)
```

```
[23]: 0.575
```

```
[24]: pd.crosstab(y_test,y_pred)
```

```
[24]: col_0    4    5    6    7
quality
3         0    1    0    0
```

```

4      1   7   2   0
5      0  98  32   0
6      0  46  76  10
7      0   3  30   9
8      0   0   1   4

```

```
[26]: f=classification_report(y_test,y_pred)
      print(f)
```

	precision	recall	f1-score	support
3	0.00	0.00	0.00	1
4	1.00	0.10	0.18	10
5	0.63	0.75	0.69	130
6	0.54	0.58	0.56	132
7	0.39	0.21	0.28	42
8	0.00	0.00	0.00	5
accuracy			0.57	320
macro avg	0.43	0.27	0.28	320
weighted avg	0.56	0.57	0.55	320

```

/usr/local/lib/python3.10/dist-packages/sklearn/metrics/_classification.py:1344:
UndefinedMetricWarning: Precision and F-score are ill-defined and being set to
0.0 in labels with no predicted samples. Use `zero_division` parameter to
control this behavior.

```

```

_warn_prf(average, modifier, msg_start, len(result))
/usr/local/lib/python3.10/dist-packages/sklearn/metrics/_classification.py:1344:
UndefinedMetricWarning: Precision and F-score are ill-defined and being set to
0.0 in labels with no predicted samples. Use `zero_division` parameter to
control this behavior.

```

```

_warn_prf(average, modifier, msg_start, len(result))
/usr/local/lib/python3.10/dist-packages/sklearn/metrics/_classification.py:1344:
UndefinedMetricWarning: Precision and F-score are ill-defined and being set to
0.0 in labels with no predicted samples. Use `zero_division` parameter to
control this behavior.

```

```
_warn_prf(average, modifier, msg_start, len(result))
```

```
[28]: probability = model.predict_proba(X_test)[: ,1]
      probability
```

```
[28]: array([0.02483477, 0.02249377, 0.01574946, 0.01134614, 0.03202448,
0.01491135, 0.00750669, 0.11853224, 0.02170618, 0.04960947,
0.01348905, 0.13794541, 0.02885185, 0.02524655, 0.02042773,
0.00977959, 0.03329229, 0.01903721, 0.00330975, 0.03185402,
0.1797074 , 0.03375518, 0.06405695, 0.01701967, 0.07199606,
0.03828524, 0.00477106, 0.08112635, 0.04058424, 0.01038622,
```


0.03132 , 0.04077697, 0.01130467, 0.04878004, 0.02851514,
0.02944521, 0.01913737, 0.06835138, 0.03249786, 0.02057124,
0.02567618, 0.02042431, 0.006988 , 0.01425204, 0.01374079,
0.03947643, 0.00374969, 0.01676378, 0.08833596, 0.07250037,
0.01094354, 0.02734811, 0.09025171, 0.01177792, 0.04743786,
0.01557644, 0.02994908, 0.09684315, 0.02995632, 0.06248571,
0.01361311, 0.02148332, 0.03028016, 0.04898658, 0.00497826,
0.01510166, 0.01183327, 0.04781076, 0.00504487, 0.01450206,
0.00975868, 0.08957871, 0.0110487 , 0.03338539, 0.02254037,
0.04057309, 0.00302335, 0.00481278, 0.02595905, 0.00514559,
0.06328622, 0.00571601, 0.02323339, 0.03167214, 0.02807292,
0.00498152, 0.01519544, 0.071952 , 0.00589653, 0.10965556,
0.00579155, 0.04130017, 0.09166548, 0.0216257 , 0.00971574,
0.01430115, 0.02018015, 0.02031068, 0.15743163, 0.02627518,
0.1286683 , 0.03879046, 0.01325419, 0.10345292, 0.04510503,
0.01106421, 0.01385973, 0.00778097, 0.03758517, 0.03038102,
0.00631245, 0.00592356, 0.00255582, 0.0085207 , 0.02894913,
0.01094736, 0.06563413, 0.01373076, 0.05589608, 0.01637058,
0.01182063, 0.0395007 , 0.01413268, 0.02768522, 0.0678964 ,
0.04112568, 0.00826422, 0.02335786, 0.03401157, 0.02896179,
0.0110487 , 0.11028868, 0.03168882, 0.00694686, 0.07250037,
0.06947725, 0.03621314, 0.00512094, 0.01887089, 0.0232298 ,
0.01022772, 0.00711342, 0.00791129, 0.05265072, 0.01430848,
0.03762969, 0.05080835, 0.13812178, 0.00655973, 0.01417962,
0.01731361, 0.02659349, 0.03403913, 0.01742294, 0.0110487 ,
0.00525754, 0.05574201, 0.01837955, 0.02863305, 0.02180492,
0.00855826, 0.01296316, 0.00719941, 0.02473188, 0.02877727,
0.01831045, 0.11902339, 0.03350339, 0.00606012, 0.01003063,
0.01231073, 0.07909799, 0.01248025, 0.00938227, 0.00922779,
0.00984898, 0.00563226, 0.01627412, 0.01257727, 0.26411334,
0.09242519, 0.00722218, 0.00895868, 0.0108839 , 0.09369835,
0.00256229, 0.14726728, 0.013673 , 0.00316726, 0.00938227,
0.05751076, 0.00648566, 0.01672627, 0.00407532, 0.01105994,
0.04465489, 0.13261131, 0.01265878, 0.02343221, 0.01106421,
0.04312645, 0.00825491, 0.02558483, 0.04303992, 0.00167707,
0.06365584, 0.02879044, 0.02108765, 0.02017862, 0.00209676,
0.04510175, 0.00560467, 0.01235811, 0.00154654, 0.01053313,
0.03851546, 0.07504775, 0.01853088, 0.04535837, 0.03899463,
0.0425178 , 0.00846753, 0.0104536 , 0.00717051, 0.23341551,
0.04884486, 0.04044068, 0.00566036, 0.00527932, 0.10618048,
0.00553059, 0.03591392, 0.01170368, 0.03609803, 0.00592857,
0.01106421, 0.06540582, 0.02807292, 0.06035659, 0.0120259 ,
0.04567327, 0.04198665, 0.06234175, 0.0061261 , 0.03219853,
0.02733513, 0.02068696, 0.04479066, 0.00227449, 0.00794547,
0.00523833, 0.0414681 , 0.04079487, 0.02742506, 0.35729187,
0.02987603, 0.02290243, 0.03998963, 0.01302718, 0.00323973,
0.01024649, 0.08956821, 0.01627412, 0.00917579, 0.04400462,

```
0.00749356, 0.03743595, 0.01220615, 0.00907912, 0.00546852,  
0.0184154 , 0.00756899, 0.01566495, 0.01252444, 0.02138715,  
0.04502192, 0.01317294, 0.00726284, 0.00533177, 0.05413836,  
0.00522663, 0.12843439, 0.0386677 , 0.02689816, 0.01919488,  
0.02910286, 0.03470307, 0.00711342, 0.03192086, 0.21076717,  
0.0268413 , 0.01016235, 0.00992186, 0.01227231, 0.01608297,  
0.00649988, 0.04338943, 0.00832624, 0.0088752 , 0.00587206,  
0.00899913, 0.02178773, 0.00842123, 0.01201218, 0.04519739,  
0.05589983, 0.01363781, 0.0050594 , 0.03334108, 0.00328764,  
0.03285411, 0.03042666, 0.02367536, 0.04470886, 0.04878004,  
0.0542435 , 0.02056979, 0.04453912, 0.00576608, 0.05107761])
```

```
[29]: model.predict([[7.4, 0.700, 0.00, 1.9, 0.076, 11.0, 34.0, 0.99780, 3.51, 0.56, ↪  
↪9.4]])
```

```
[29]: array([5])
```

```
[ ]:
```